





Elasticsearch practices at eBay

Pei Wang

Aug 2018



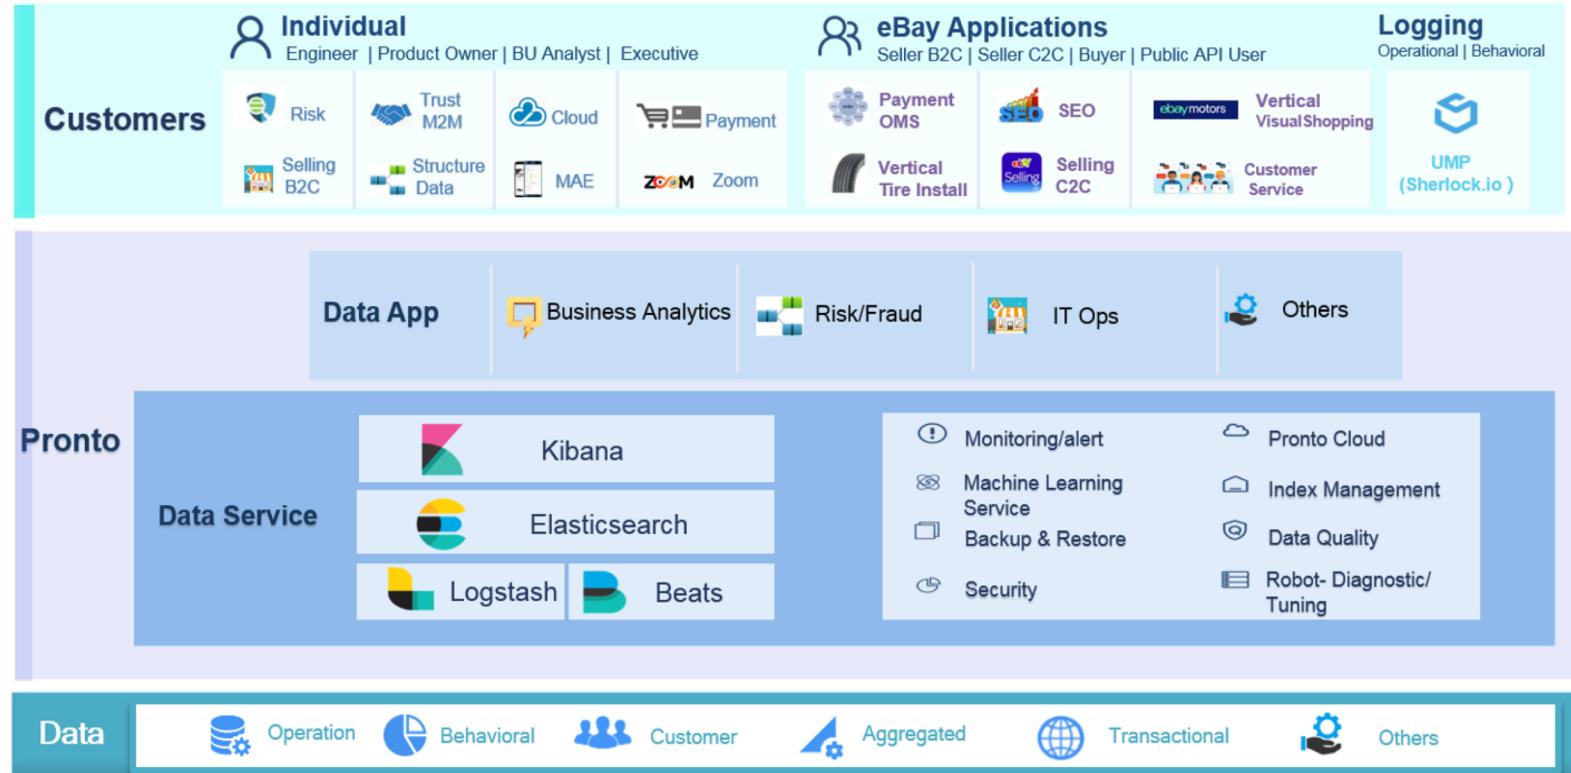
Agenda

- Elasticsearch in eBay tech stack
- Performance tuning
- Index management

1

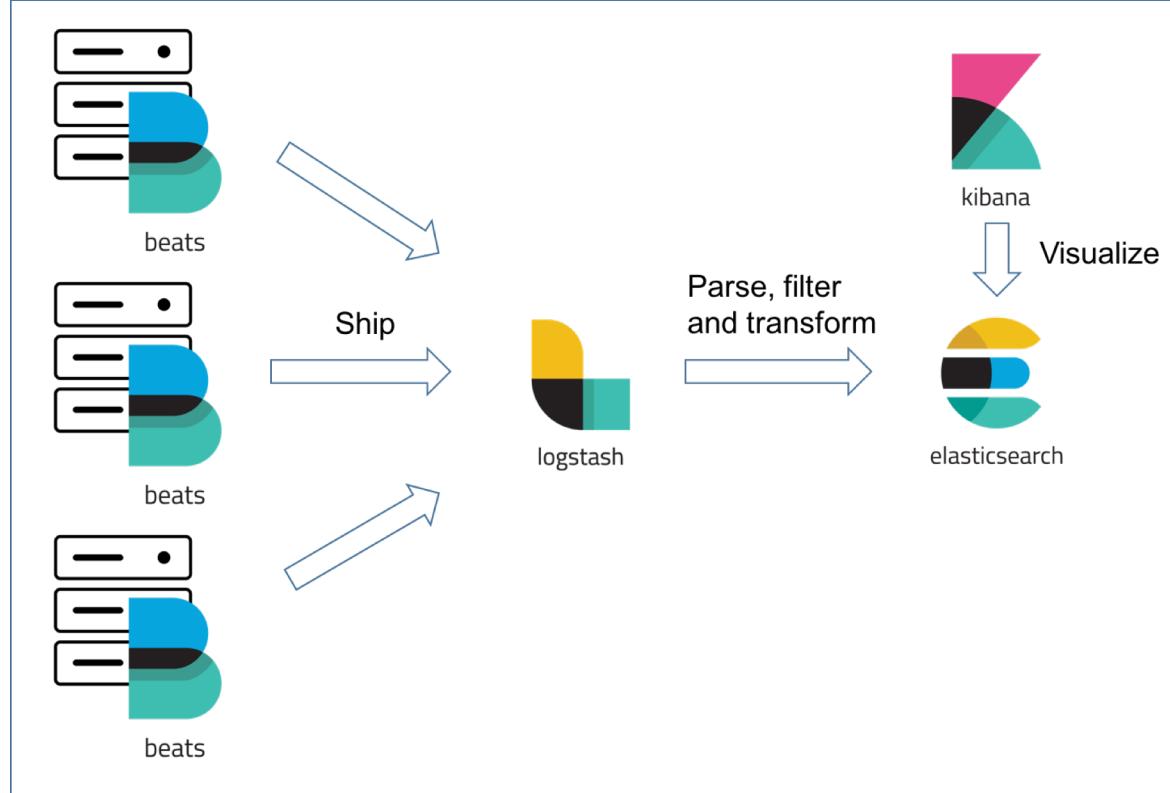
Elasticsearch in eBay tech stack

Elasticsearch in eBay tech stack



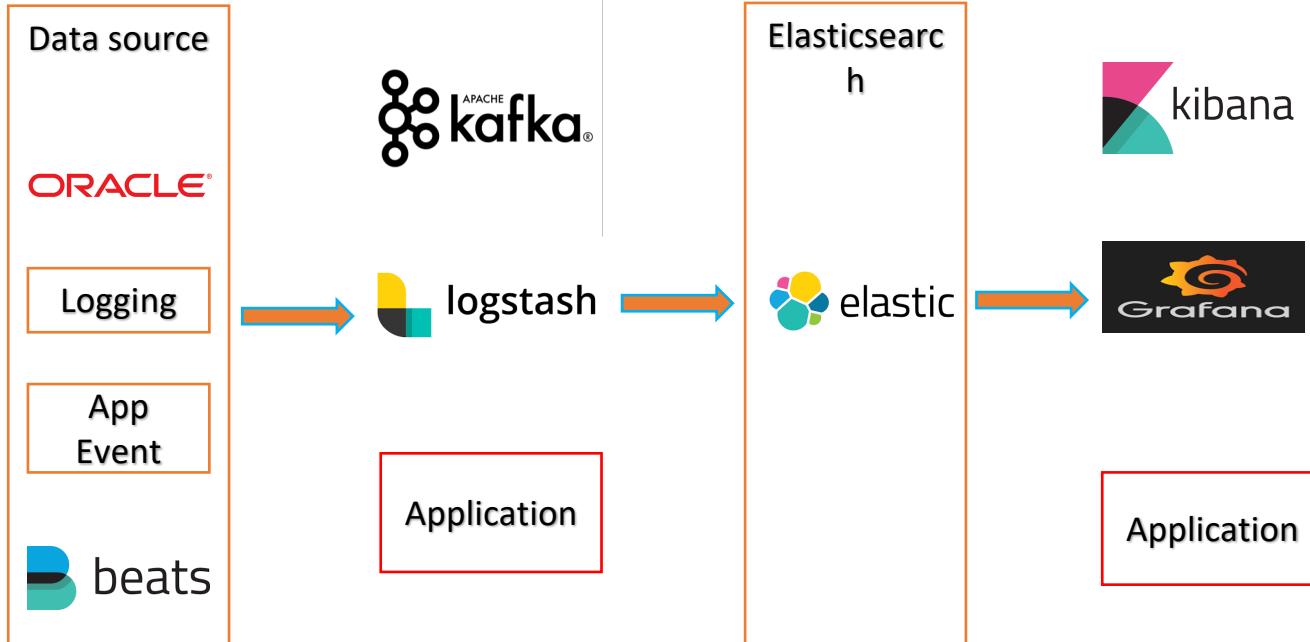


ELK + Beats





Elasticsearch in eBay tech stack





Agenda

- Elasticsearch in eBay tech stack
- Performance tuning
- Index management

2

Aggregation



Optimize Index Design

Let us consider such scenario

- there is an index which have one billion social media messages
- we have queries like below.

```
GET _search
{
  "query": {
    "bool": {
      "must": [
        { "match": { "title": "Search" } }
      ],
      "filter": [
        { "term": { "region": "US" } },
        { "range": { "publish_date": { "gte": "2017-01-01" } } },
        { "match": { "author_id": 10086 } }
      ]
    }
  }
}
```



Optimize Index Design

- **Organize your index by date**
 - Cases: Logging / Monitoring / Event
- **Organize your index by field, like region**
 - Cases: Query with **enumerable** filter field
- **Use routing key**
 - Distribute docs with same routing key to same shard
 - Cases: Query with **non-enumerable** filter field

```
GET _search
{
  "query": {
    "bool": {
      "must": [
        { "match": { "title": "Search" } }
      ],
      "filter": [
        { "term": { "region": "US" } },
        { "range": { "publish_date": { "gte": "2017-01-01" } } },
        { "match": { "author_id": 10086 } }
      ]
    }
  }
}
```



Optimize Index Design

- **Set mapping explicitly**
 - The default mapping may not fit your case
- **Make shards distributed evenly across nodes**
 - Nodes have more shards than others may become the bottle neck
- **Avoid imbalanced sharding if docs are indexed with routing key or user-defined ID**
 - User-defined ID should be random enough.
 - Imbalanced routing key distribution could cause imbalanced sharding.

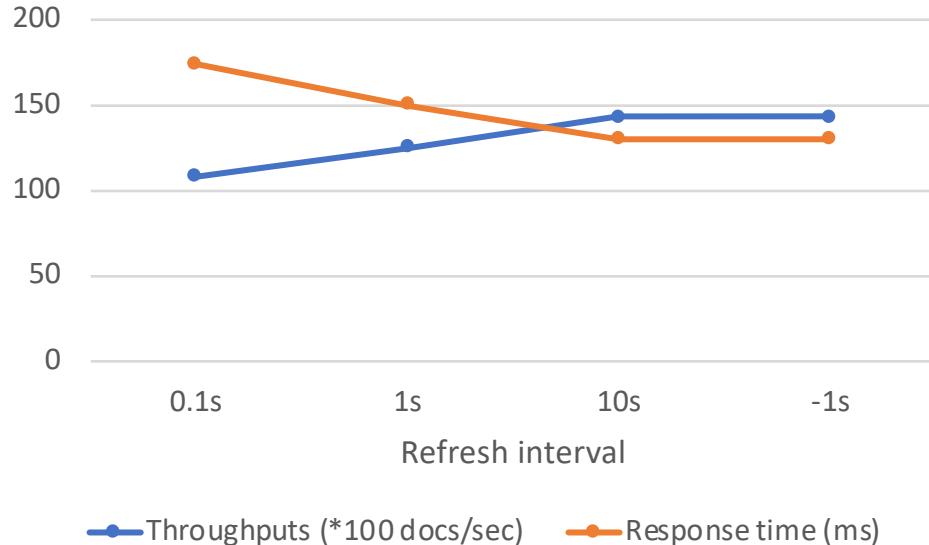
atest				
shards: 4 * 1 docs: 148,888 size: 369.61MB				
node-1	10.64.206.171	heap	disk	cpu
2				
node-2	10.64.206.173	heap	disk	cpu
0				
node-3	10.64.206.174	heap	disk	cpu
1	3			





Index performance tuning

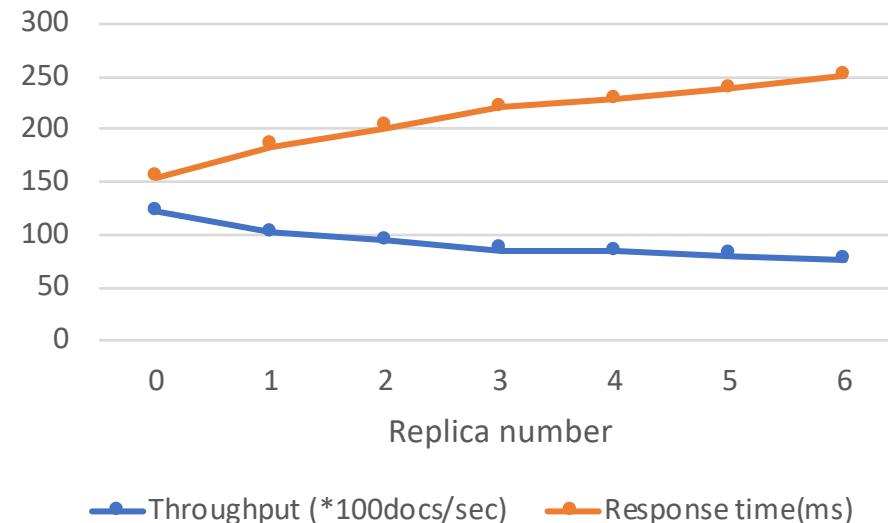
- **Increase shard number if node count > (1+ replica) * shard number**
 - Scale out, distribute data into more nodes
- **Increase refresh interval**
 - ElasticSearch create a segment every time refresh event happen. Increase interval would reduce segment count and merge cost.
 - Documents are not available for search until refresh





Index performance tuning

- **Use auto generated IDs if possible**
 - ElasticSearch auto generate ID algorithm can reduce the duplicate ID check and version check cost.
- **Reduce replica number**
 - ElasticSearch need to write primary shard and all replica shards for every index request
 - Replica number should not be 0, otherwise you will have data loss risk

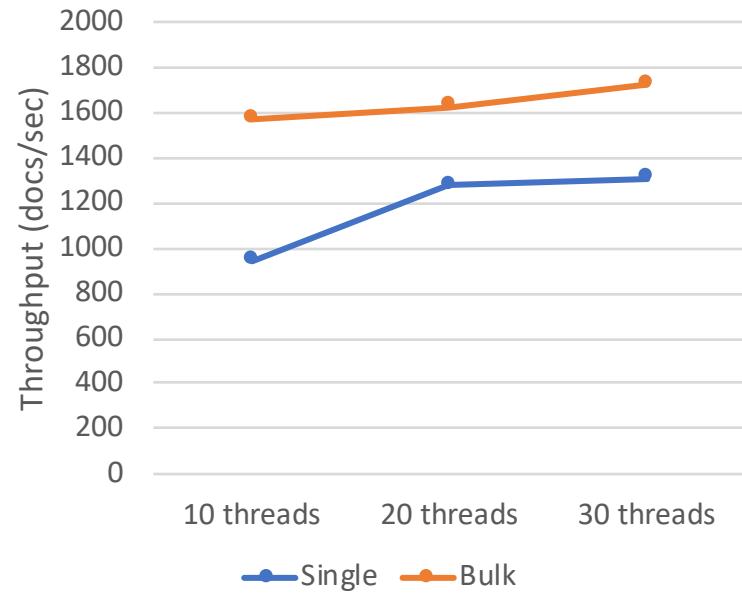




Index performance tuning

Client side

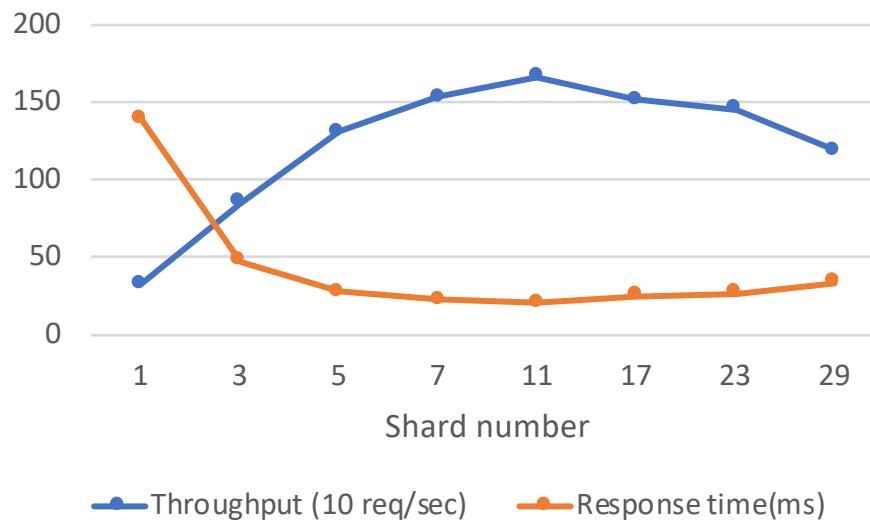
- **Use bulk request**
- **Use multiple threads/workers**





Search performance tuning

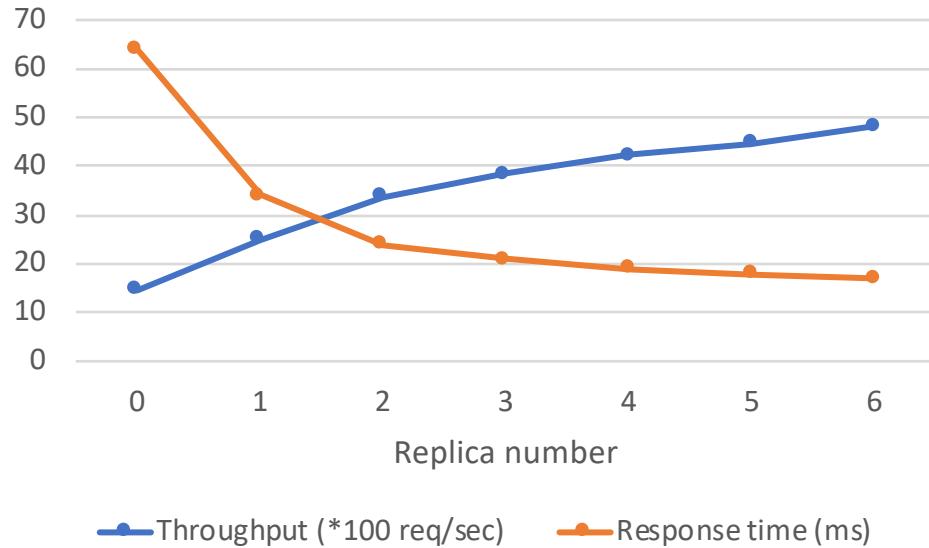
- **Choose suitable shard number**
 - Too small shard number will make search unable to scale out.
 - Too big shard number will hurt performance too.
 - Shard size should not exceed 30-50GB
 - Notes, in this test, every shard has an exclusive node.





Search performance tuning

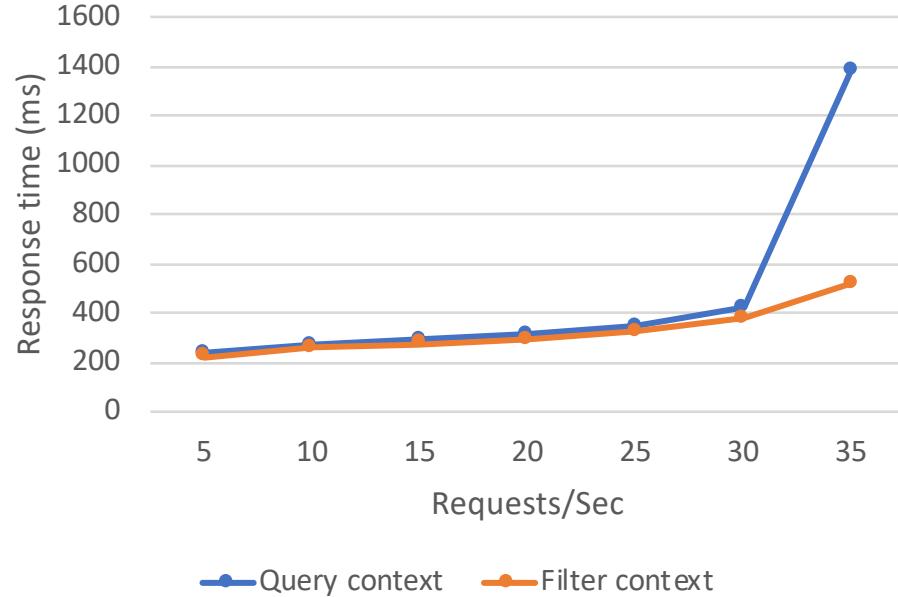
- **Increase replica number**
 - Search can be performed on either a primary or replica shard.
 - Will decrease indexing performance
 - Notes, in this test, every shard has an exclusive node.





Search performance tuning

- **Use filter context instead of query context if possible**
 - ElasticSearch do not need to calculate relevancy score for filter context
 - Filtered result can be cached





Search performance tuning

- **Node query cache**

- Only cache queries used in filter context.
- ElasticSearch used bit set mechanism to cache filter results.
- ElasticSearch only enable node query cache for segments have more than 10000 or 3% documents, whichever is larger.

```
GET index_name/_stats?filter_path=indices.**.query_cache
{
  "indices": {
    "index_name": {
      "primaries": {
        "query_cache": {
          "memory_size_in_bytes": 46004616,
          "total_count": 1588886,
          "hit_count": 515001,
          "miss_count": 1073885,
          "cache_size": 630,
          "cache_count": 630,
          "evictions": 0
        }
      },
      "total": {
        "query_cache": {
          "memory_size_in_bytes": 46004616,
          "total_count": 1588886,
          "hit_count": 515001,
          "miss_count": 1073885,
          "cache_size": 630,
          "cache_count": 630,
          "evictions": 0
        }
      }
    }
  }
}
```



Search performance tuning

- **Shard request cache**

- Only cache request size:0, like aggregate, suggestions or hit.totals.
- Use payload body as cache key, so that the payload must be the same, even the JSON body order.
- Do not use now or new Date() to build request body since it will make payload change every time, round your date time.
- Cache are invalidated when refresh happen and the data in the shard has actually changed.

```
GET index_name/_stats?filter_path=indices.*.request_cache
{
  "indices": {
    "index_name": {
      "primaries": {
        "request_cache": {
          "memory_size_in_bytes": 0,
          "evictions": 0,
          "hit_count": 541,
          "miss_count": 514098
        }
      },
      "total": {
        "request_cache": {
          "memory_size_in_bytes": 0,
          "evictions": 0,
          "hit_count": 982,
          "miss_count": 947321
        }
      }
    }
  }
}
```



Search performance tuning

- **Retrieve only necessary fields**
 - Use “_source” or “stored_fileds” to let ElasticSearch only return necessary fields
- **Reduce documents count in response if possible**

```
GET /perf_test_result/_search
{
  "size": 5,
  "_source": [
    "job_name",
    "create_at"
  ]
}

{
  "hits": {
    "total": 12338,
    "max_score": 1,
    "hits": [
      {
        "_index": "perf_test_result",
        "_type": "REQUEST",
        "_id": "AVwv0bQplyIgCso4VHVs",
        "_score": 1,
        "_source": {
          "job_name": "single_shard_ingest_test",
          "create_at": "2017-05-22T08:15:29.703Z"
        }
      },
      ...
    ]
}
```



Search performance tuning

- Sort by “_doc” explicitly if do not care about the document order in response
 - ElasticSearch use the “_score” field to sort by score as default.
 - Use “sort”: “_doc” to let ElasticSearch return hits by index order
 - Especially helps when scrolling

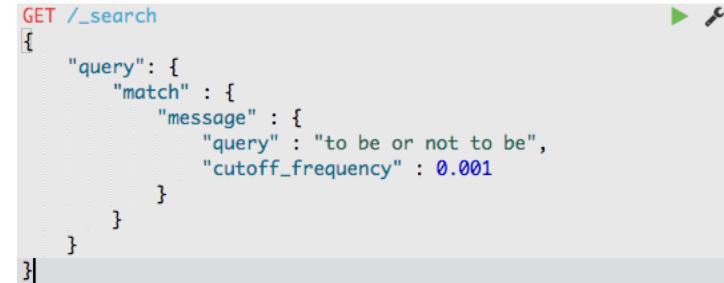
```
GET /perf_test_result/_search
{
  "sort": ["_doc"]
}
```





Search performance tuning

- **Avoid searching stop words**
 - Stop words like "a" and "the" may cause the query hit results count to explode.
 - Use stop word filter
 - Refine query, use "the AND fox" if you really meant to search "the" word.
 - If some words are frequently used in your documents but not in the default list, use cutoff_frequency to specify which terms are stop words.



```
GET /_search
{
  "query": {
    "match": {
      "message": {
        "query" : "to be or not to be",
        "cutoff_frequency" : 0.001
      }
    }
  }
}
```



Search performance tuning

- **Avoid using script query to calculate hits in flight.**
 - Script query like below is quite time-consuming.
 - Should consider add extra fields in indexing phase if you have a lot of script query

```
{  
  "query": {  
    "script" : {  
      "script" : {  
        "inline": "doc['job_name.keyword'].value.startsWith('ebay')"  
      }  
    }  
  }  
}
```



```
{  
  "query": {  
    "match": {  
      "job_name_prefix.keyword": "ebay"  
    }  
  }  
}
```





Search performance tuning

- **Avoid wildcard query**
 - wildcard query like below is quite time-consuming, like script query, especially for leading wildcard query like “*ebay”
 - Should consider add extra fields in indexing phase if you have a lot of wildcard query

```
{  
    "query": {  
        "wildcard": {  
            "job_name.keyword": "ebay*"  
        }  
    }  
}
```



```
{  
    "query": {  
        "match": {  
            "job_name_prefix.keyword": "ebay"  
        }  
    }  
}
```





Agenda

- Elasticsearch in eBay tech stack
- Performance tuning
- Index management

3

Index management



Index Lifecycle Management Service

What it can do

- High availability index lifecycle management with user-friendly web portal
- One-stop management portal for multiple clusters and multiple tenants

Why we build it

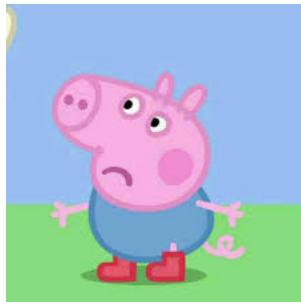
- Help customers manage ES data through its full lifecycle
- Let important business applications survive from peak days
- Avoid improper ES usage.



Bad design

All data in one index

10 million



As time increase



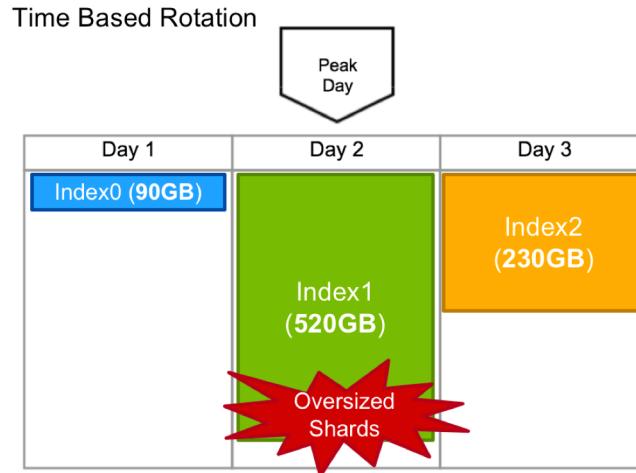
1 billion





Time Based Creation

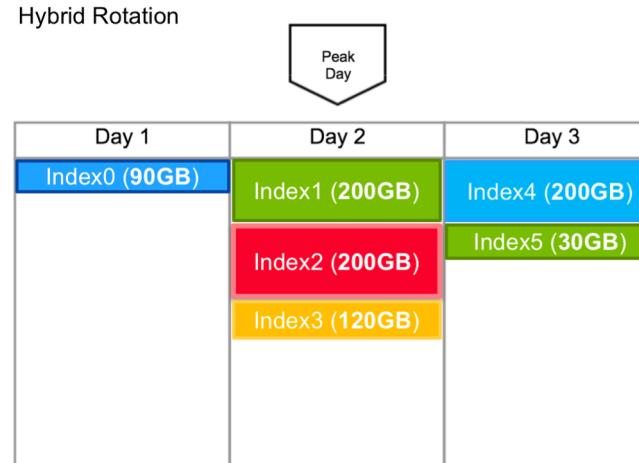
- Create new index according to timestamp automatically
- Distribute document base on timestamp
- Delete expired indices automatically





Hybrid Rotation

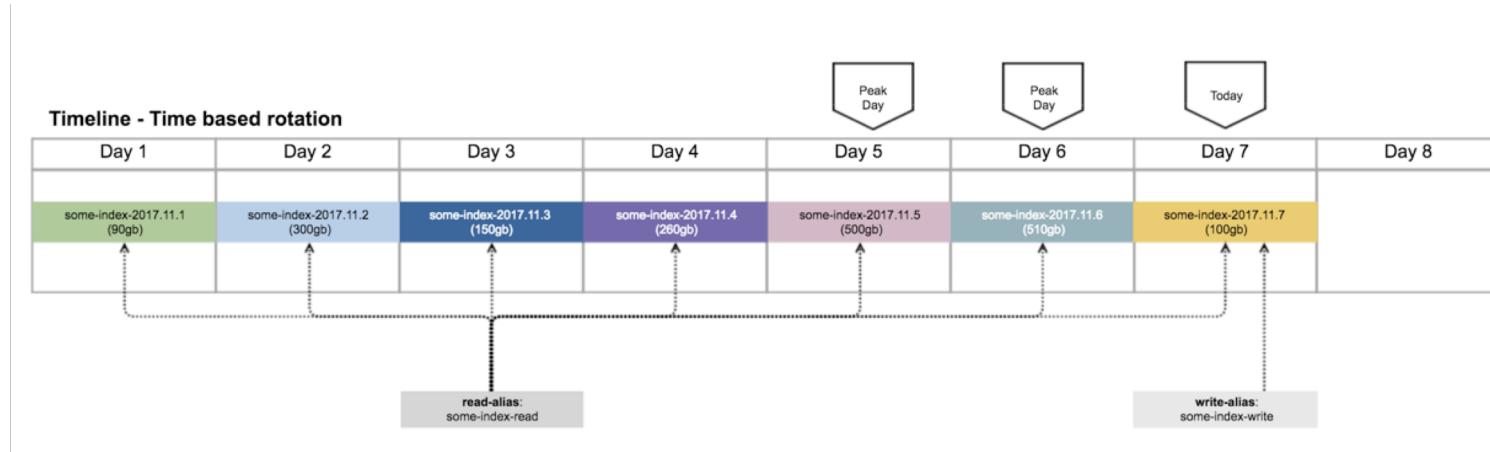
- Besides time rotation, Hybrid solution take disk space into consideration
- Create and manage indices automatically





Index alias

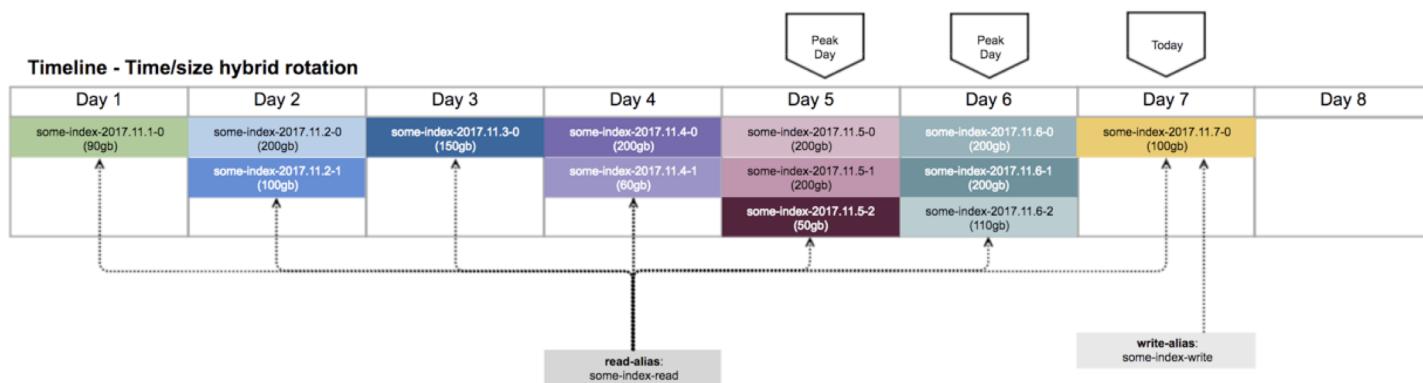
- Minimize customer's effort
- Use Index alias to route user request to correct indices





Index alias

- Minimize customer's effort
- Use Index alias to route user request to correct indices





Rule Overview

Index Management

test

DASHBOARD

CLUSTERS

RULES

JOB

Rule Add

Cluster Name: Please select cluster

Index Prefix: require not to be empty

Timezone Setting: -0700

Alias Patterns: read_all write

Last 5 day

Index Rotation Strategies: create close purge setting_change

Creation Type: hybrid

Index Creation Interval: 1 Please select time unit

Max size per shard: 30 GB

Default ES settings: require json format

Close indices older than: 0 Please select time unit

Purge indices older than: 0 Please select time unit

Setting change rules:

Change settings to below after 0 Please select time unit

Submit Cancel



Action: Index Creation

Index Rotation Strategies: **create** close purge setting_change

Creation Type: hybrid

Index Creation Interval: 1 Please select time unit

Max size per shard: 30 GB

Default ES settings: require json format

Submit **Cancel**



Action: Index Close and Delete

Index Rotation Strategies: create close purge setting_change

Close indices older than: Please select time unit ▾

- day
- week
- month

Index Rotation Strategies: create close purge setting_change

Purge indices older than: Please select time unit ▾

- day
- week
- month



Alias Rotation

Alias Patterns:

read_all write +

Last	5	Please select time unit ^	×
Last		day	×
Last		week	×
Last		month	×



Action: Index Settings

Index Rotation Strategies: create close purge setting_change

Setting change rules:

Change settings to below after week

```
{  
  "index": {  
    "routing": {  
      "allocation": {  
        "require": {  
          "box_type": "warm"  
        }  
      }  
    }  
  }  
}
```

Submit **Cancel**



Feature Summary

Business Specifications

- Index Creation
 - Condition :Time, Size, Shard-Size
 - Supported time unit : hour, day, week, month
 - Supported size unit : GB
- Index Close
- Index Deletion
- Index Setting Change
 - Condition :Time
 - Supported time unit : hour, day, week, month
 - Use case: hot-warm, replica number, etc.
- Alias Rotation

Rule Add

Cluster Name: Please select cluster

Index Prefix: require not to be empty

Timezone Setting: -0700

Alias Patterns: read_all write

Index Rotation create close purge setting_change

Strategies:

Creation Type: hybrid

Index Creation Interval: 1 Please select time unit

Max size per shard: 30 GB

Default ES settings: require json format

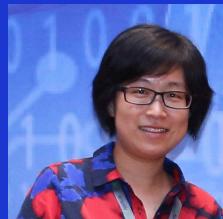
Submit Cancel

test

Thanks



Wang, Pei



Dai, Catherine



Hu, Hongsi